



Cairo University
Faculty of Engineering Credit Hours System

CMPS401
Advanced Database Systems
Project Proposal

Prepared by:

- Abdelrahman Mohamed Ezzat 1190158
- Abdelrahman Ibrahim Mohamed 1190493
- Amr Ahmed Abdelzaher 1190074
- Joseph Ameer Aziz 1190052

DiskANN-based index

We are going to implement an approximate nearest neighbour algorithm based on DiskANN[1], which uses Vamana data structure. This is a graph-based index that uses Greedy Search for retrieving the top “**k**” approximate nearest neighbours to a given vector embedding. The algorithm was modified in multiple papers but we are referencing the original paper published in 2019.

The algorithm optimizes over brute force comparison that it compares the query with much fewer nodes.

The algorithm at the very beginning constructs a random R -regular directed graph, which means that the out-degree of each node is at most R . Each node represents an embedding and the edges represent how close the node is to its neighbours.

Graph Construction

1. We initialize a graph where each node has R randomly chosen out-neighbours.
2. We chose the vertex s (the medoid of the dataset) which will be the starting node for the search algorithm.
3. We iterate on all the points in a random order and then update the graph to make the greedy search converge faster.
4. We run Greedy Search from every point with ($k = 1$) and then we apply a pruning algorithm to maintain the R out-degree property.

Algorithm 1: GreedySearch(s, x_q, k, L)

Data: Graph G with start node s , query x_q , result size k , search list size $L \geq k$

Result: Result set \mathcal{L} containing k -approx NNs, and a set \mathcal{V} containing all the visited nodes

```
begin
  initialize sets  $\mathcal{L} \leftarrow \{s\}$  and  $\mathcal{V} \leftarrow \emptyset$ 
  while  $\mathcal{L} \setminus \mathcal{V} \neq \emptyset$  do
    let  $p^* \leftarrow \arg \min_{p \in \mathcal{L} \setminus \mathcal{V}} \|x_p - x_q\|$ 
    update  $\mathcal{L} \leftarrow \mathcal{L} \cup N_{\text{out}}(p^*)$  and
       $\mathcal{V} \leftarrow \mathcal{V} \cup \{p^*\}$ 
    if  $|\mathcal{L}| > L$  then
      update  $\mathcal{L}$  to retain closest  $L$ 
        points to  $x_q$ 
  return [closest  $k$  points from  $\mathcal{L}$ ;  $\mathcal{V}$ ]
```

Algorithm 2: RobustPrune($p, \mathcal{V}, \alpha, R$)

Data: Graph G , point $p \in P$, candidate set \mathcal{V} , distance threshold $\alpha \geq 1$, degree bound R

Result: G is modified by setting at most R new out-neighbors for p

```
begin
   $\mathcal{V} \leftarrow (\mathcal{V} \cup N_{\text{out}}(p)) \setminus \{p\}$ 
   $N_{\text{out}}(p) \leftarrow \emptyset$ 
  while  $\mathcal{V} \neq \emptyset$  do
     $p^* \leftarrow \arg \min_{p' \in \mathcal{V}} d(p, p')$ 
     $N_{\text{out}}(p) \leftarrow N_{\text{out}}(p) \cup \{p^*\}$ 
    if  $|N_{\text{out}}(p)| = R$  then
      break
    for  $p' \in \mathcal{V}$  do
      if  $\alpha \cdot d(p^*, p') \leq d(p, p')$  then
        remove  $p'$  from  $\mathcal{V}$ 
```

K Top ANNs Retrieval

We use Greedy Search with the given k after constructing the graph correctly. We also use cosine similarity for distance comparison.

References

[1] DiskANN: Fast Accurate Billion-point Nearest Neighbor Search on a Single Node:
https://proceedings.neurips.cc/paper_files/paper/2019/file/09853c7fb1d3f8ee67a61b6bf4a7f8e6-Paper.pdf